

Deep Learning Project: Credit Card Fraud Detection

Objective

The main objective of this analysis is to develop a deep learning model to accurately detect fraudulent credit card transactions. This task focuses on binary classification using supervised learning techniques, with an emphasis on addressing the significant class imbalance in the dataset. Insights from this analysis can help credit card companies reduce financial losses and enhance customer satisfaction by identifying fraudulent transactions efficiently.

Dataset Description

The dataset comes from (<https://www.kaggle.com/datasets/joebeachcapital/credit-card-fraud>) and contains credit card transactions made by European cardholders in September 2013. It includes a total of 284,807 transactions, of which only 492 (0.172%) are fraudulent, highlighting a significant class imbalance. Features V1 to V28 are Principal Component Analysis (PCA) transformations of the original features, with 'Time' and 'Amount' being untransformed. The 'Class' column is the target variable, where 0 represents non-fraud and 1 represents fraud.

```
In [65]: # Load the dataset
import pandas as pd
import numpy as np

file_path = r'C:\Users\gcantarella\Desktop\Deep Learning Final Project\creditcard.csv'
data = pd.read_csv(file_path)

# Summarize dataset structure
print(data.info())
# Display a preview of the dataset
print(data.head())
# Check for missing values
print(data.isnull().sum())
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Time    284807 non-null   float64
1   V1      284807 non-null   float64
2   V2      284807 non-null   float64
3   V3      284807 non-null   float64
4   V4      284807 non-null   float64
5   V5      284807 non-null   float64
6   V6      284807 non-null   float64
7   V7      284807 non-null   float64
8   V8      284807 non-null   float64
9   V9      284807 non-null   float64
10  V10     284807 non-null   float64
11  V11     284807 non-null   float64
12  V12     284807 non-null   float64
13  V13     284807 non-null   float64
14  V14     284807 non-null   float64
15  V15     284807 non-null   float64
16  V16     284807 non-null   float64
17  V17     284807 non-null   float64
18  V18     284807 non-null   float64
19  V19     284807 non-null   float64
20  V20     284807 non-null   float64
21  V21     284807 non-null   float64
22  V22     284807 non-null   float64
23  V23     284807 non-null   float64
24  V24     284807 non-null   float64
25  V25     284807 non-null   float64
26  V26     284807 non-null   float64
27  V27     284807 non-null   float64
28  V28     284807 non-null   float64
29  Amount  284807 non-null   float64
30  Class   284807 non-null   int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB

```

```

None
      Time      V1      V2      V3      V4      V5      V6      V7  \
0   0.0 -1.359807 -0.072781  2.536347  1.378155 -0.338321  0.462388  0.239599
1   0.0  1.191857  0.266151  0.166480  0.448154  0.060018 -0.082361 -0.078803
2   1.0 -1.358354 -1.340163  1.773209  0.379780 -0.503198  1.800499  0.791461
3   1.0 -0.966272 -0.185226  1.792993 -0.863291 -0.010309  1.247203  0.237609
4   2.0 -1.158233  0.877737  1.548718  0.403034 -0.407193  0.095921  0.592941

      V8      V9  ...      V21      V22      V23      V24      V25  \
0  0.098698  0.363787  ... -0.018307  0.277838 -0.110474  0.066928  0.128539
1  0.085102 -0.255425  ... -0.225775 -0.638672  0.101288 -0.339846  0.167170
2  0.247676 -1.514654  ...  0.247998  0.771679  0.909412 -0.689281 -0.327642
3  0.377436 -1.387024  ... -0.108300  0.005274 -0.190321 -1.175575  0.647376
4 -0.270533  0.817739  ... -0.009431  0.798278 -0.137458  0.141267 -0.206010

      V26      V27      V28  Amount  Class
0 -0.189115  0.133558 -0.021053  149.62    0
1  0.125895 -0.008983  0.014724   2.69    0
2 -0.139097 -0.055353 -0.059752  378.66    0
3 -0.221929  0.062723  0.061458  123.50    0
4  0.502292  0.219422  0.215153   69.99    0

```

```
[5 rows x 31 columns]
Time      0
V1        0
V2        0
V3        0
V4        0
V5        0
V6        0
V7        0
V8        0
V9        0
V10       0
V11       0
V12       0
V13       0
V14       0
V15       0
V16       0
V17       0
V18       0
V19       0
V20       0
V21       0
V22       0
V23       0
V24       0
V25       0
V26       0
V27       0
V28       0
Amount    0
Class     0
dtype: int64
```

Exploratory Data Analysis (EDA)

```
In [30]: import matplotlib.pyplot as plt
import seaborn as sns

# Check the class distribution
class_distribution = data['Class'].value_counts(normalize=True) * 100
print(f"Class distribution: \n{class_distribution}")

# Visualize the distribution of the target variable
plt.figure(figsize=(6, 4))
sns.barplot(x=class_distribution.index, y=class_distribution.values, palette="vi
plt.title('Class Distribution (Fraud vs Non-Fraud)')
plt.ylabel('Proportion')
plt.xlabel('Class (0: Non-Fraud, 1: Fraud)')
plt.xticks([0, 1], ['Non-Fraud', 'Fraud'])
plt.legend([], [], frameon=False) # Removes redundant Legend
plt.show()


# Describe the dataset to understand feature statistics
data_description = data.describe()

class_distribution
data_description.iloc[:, :8]
```

```

Class distribution:
Class
0    99.827251
1     0.172749
Name: proportion, dtype: float64

```



Out[30]:

	Time	V1	V2	V3	V4	
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.8
mean	94813.859575	1.168375e-15	3.416908e-16	-1.379537e-15	2.074095e-15	9.6
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00	1.3
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	-1.1
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01	-6.5
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02	-5.4
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01	6.1
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01	3.4

Data Preprocessing

The following steps were undertaken for preprocessing the dataset:

- Normalized 'Time' and 'Amount' using StandardScaler to bring all features onto a similar scale.
- Split the dataset into training (80%) and testing (20%) sets, with stratified sampling to maintain class balance.

```
In [32]: # Normalize 'Amount' and 'Time' features
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
data[['Time', 'Amount']] = scaler.fit_transform(data[['Time', 'Amount']])

# Split data into features and target
X = data.drop(columns=['Class'])
y = data['Class']

# Train-test split
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_

# Confirm shapes of the datasets
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
Out[32]: ((227845, 30), (56962, 30), (227845,), (56962,))
```

Model Development

Three deep learning models were developed and evaluated for this task:

1. A simple neural network with two dense layers and a dropout layer.
2. A more complex neural network with additional layers and nodes.
3. A CNN will be implemented.

Simple Neural Network

```
In [52]: # Import necessary libraries
import tensorflow as tf
from tensorflow.keras import Sequential, Input
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.callbacks import EarlyStopping
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

# Preprocessing
# Assuming `data` is a pandas DataFrame containing the dataset
# Separate features and target variable
X = data.drop(columns=['Class'])
y = data['Class']

# Normalize 'Time' and 'Amount' features
scaler = StandardScaler()
X[['Time', 'Amount']] = scaler.fit_transform(X[['Time', 'Amount']])

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# Define the simple neural network model
def build_simple_model(input_shape):
```

```


model = Sequential([
    Input(shape=(input_shape,)), # Use Input Layer to specify the input sha
    Dense(16, activation='relu'),
    Dropout(0.5),
    Dense(8, activation='relu'),
    Dense(1, activation='sigmoid') # Output layer for binary classification
])
model.compile(
    optimizer='adam',
    loss='binary_crossentropy',
    metrics=['AUC']
)
return model


# Build and compile the model
simple_model = build_simple_model(X_train.shape[1])


# Early stopping to prevent overfitting
early_stopping = EarlyStopping(
    monitor='val_loss', # Monitor validation loss
    patience=8, # Number of epochs with no improvement
    restore_best_weights=True # Restore the best model weights
)


# Train the model
history = simple_model.fit(
    X_train, y_train,
    validation_split=0.2, # Use 20% of the training data for validation
    epochs=20, # Train for a maximum of 20 epochs
    batch_size=512, # Use a batch size of 512
    callbacks=[early_stopping], # Early stopping callback
    verbose=1 # Print training progress
)


```


Epoch 1/20
357/357  2s 2ms/step - AUC: 0.4319 - loss: 0.3550 - val_AUC: 0.8596 - val_loss: 0.0118


Epoch 2/20
357/357  1s 2ms/step - AUC: 0.8167 - loss: 0.0262 - val_AUC: 0.9150 - val_loss: 0.0053


Epoch 3/20
357/357  1s 2ms/step - AUC: 0.9097 - loss: 0.0129 - val_AUC: 0.9154 - val_loss: 0.0047


Epoch 4/20
357/357  1s 2ms/step - AUC: 0.9097 - loss: 0.0091 - val_AUC: 0.9215 - val_loss: 0.0045


Epoch 5/20
357/357  1s 2ms/step - AUC: 0.9300 - loss: 0.0071 - val_AUC: 0.9274 - val_loss: 0.0043


Epoch 6/20
357/357  1s 2ms/step - AUC: 0.9344 - loss: 0.0063 - val_AUC: 0.9333 - val_loss: 0.0042


Epoch 7/20
357/357  1s 2ms/step - AUC: 0.9244 - loss: 0.0061 - val_AUC: 0.9333 - val_loss: 0.0042


Epoch 8/20
357/357  1s 2ms/step - AUC: 0.9455 - loss: 0.0062 - val_AUC: 0.9334 - val_loss: 0.0042


Epoch 9/20
357/357  1s 2ms/step - AUC: 0.9249 - loss: 0.0059 - val_AUC: 0.9334 - val_loss: 0.0041


Epoch 10/20
357/357  1s 2ms/step - AUC: 0.9428 - loss: 0.0058 - val_AUC: 0.9274 - val_loss: 0.0041


Epoch 11/20
357/357  1s 1ms/step - AUC: 0.9389 - loss: 0.0055 - val_AUC: 0.9273 - val_loss: 0.0043


Epoch 12/20
357/357  1s 2ms/step - AUC: 0.9371 - loss: 0.0050 - val_AUC: 0.9273 - val_loss: 0.0041


Epoch 13/20
357/357  1s 2ms/step - AUC: 0.9524 - loss: 0.0045 - val_AUC: 0.9273 - val_loss: 0.0039


Epoch 14/20
357/357  1s 2ms/step - AUC: 0.9525 - loss: 0.0051 - val_AUC: 0.9273 - val_loss: 0.0040


Epoch 15/20
357/357  1s 2ms/step - AUC: 0.9206 - loss: 0.0067 - val_AUC: 0.9333 - val_loss: 0.0041

Epoch 16/20
357/357  1s 2ms/step - AUC: 0.9281 - loss: 0.0050 - val_AUC: 0.9333 - val_loss: 0.0038

Epoch 17/20
357/357  1s 1ms/step - AUC: 0.9577 - loss: 0.0044 - val_AUC: 0.9332 - val_loss: 0.0038

Epoch 18/20
357/357  1s 1ms/step - AUC: 0.9458 - loss: 0.0051 - val_AUC: 0.9333 - val_loss: 0.0038

Epoch 19/20
357/357  1s 2ms/step - AUC: 0.9582 - loss: 0.0042 - val_AUC: 0.9333 - val_loss: 0.0038

Epoch 20/20
357/357  1s 2ms/step - AUC: 0.9488 - loss: 0.0050 - val_AUC: 0.9334 - val_loss: 0.0041

More complex Neural Network


















```
In [41]: # Define a more complex neural network
def build_complex_model(input_shape):
    model = Sequential([
        Input(shape=(input_shape,)),
        Dense(64, activation='relu'),
        Dropout(0.4),
        Dense(32, activation='relu'),
        Dropout(0.3),
        Dense(16, activation='relu'),
        Dense(1, activation='sigmoid')
    ])
    model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['AUC'])
    return model

# Build and compile the complex model
complex_model = build_complex_model(X_train.shape[1])

# Train the complex model
history_complex = complex_model.fit(
    X_train, y_train,
    validation_split=0.2,
    epochs=20,
    batch_size=512,
    callbacks=[tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=5,
    verbose=1
    )
    ]
)
```



```

Epoch 1/20
357/357  2s 3ms/step - AUC: 0.4348 - loss: 0.1570 - val_AUC:
0.9330 - val_loss: 0.0044
Epoch 2/20
357/357  1s 3ms/step - AUC: 0.8999 - loss: 0.0064 - val_AUC:
0.9272 - val_loss: 0.0038
Epoch 3/20
357/357  1s 3ms/step - AUC: 0.9315 - loss: 0.0056 - val_AUC:
0.9333 - val_loss: 0.0037
Epoch 4/20
357/357  1s 3ms/step - AUC: 0.9454 - loss: 0.0038 - val_AUC:
0.9333 - val_loss: 0.0035
Epoch 5/20
357/357  1s 3ms/step - AUC: 0.9340 - loss: 0.0041 - val_AUC:
0.9334 - val_loss: 0.0034
Epoch 6/20
357/357  1s 3ms/step - AUC: 0.9557 - loss: 0.0035 - val_AUC:
0.9393 - val_loss: 0.0033
Epoch 7/20
357/357  1s 3ms/step - AUC: 0.9481 - loss: 0.0034 - val_AUC:
0.9334 - val_loss: 0.0033
Epoch 8/20
357/357  1s 2ms/step - AUC: 0.9549 - loss: 0.0035 - val_AUC:
0.9393 - val_loss: 0.0033
Epoch 9/20
357/357  1s 3ms/step - AUC: 0.9475 - loss: 0.0028 - val_AUC:
0.9334 - val_loss: 0.0033
Epoch 10/20
357/357  1s 2ms/step - AUC: 0.9462 - loss: 0.0032 - val_AUC:
0.9334 - val_loss: 0.0032
Epoch 11/20
357/357  1s 3ms/step - AUC: 0.9618 - loss: 0.0029 - val_AUC:
0.9452 - val_loss: 0.0031
Epoch 12/20
357/357  1s 3ms/step - AUC: 0.9606 - loss: 0.0030 - val_AUC:
0.9392 - val_loss: 0.0030
Epoch 13/20
357/357  1s 3ms/step - AUC: 0.9451 - loss: 0.0029 - val_AUC:
0.9453 - val_loss: 0.0031
Epoch 14/20
357/357  1s 3ms/step - AUC: 0.9609 - loss: 0.0032 - val_AUC:
0.9393 - val_loss: 0.0032
Epoch 15/20
357/357  1s 3ms/step - AUC: 0.9638 - loss: 0.0031 - val_AUC:
0.9334 - val_loss: 0.0033
Epoch 16/20
357/357  1s 3ms/step - AUC: 0.9615 - loss: 0.0025 - val_AUC:
0.9334 - val_loss: 0.0033
Epoch 17/20
357/357  1s 3ms/step - AUC: 0.9473 - loss: 0.0025 - val_AUC:
0.9394 - val_loss: 0.0033

```

CNN Implementation

We can try using a **Convolutional Neural Network (CNN)**. CNNs are particularly effective at identifying spatial patterns and interactions between features. To apply a CNN to tabular data, we need to reshape the data into a format that mimics an image structure.

Here's how to proceed:

1. Reshape the Input Data
2. Convert the flat feature vectors into a 2D matrix. For example, the 28 PCA features + 2 others could be reshaped into a 6x5 or similar grid.
3. Build a CNN Architecture
4. Add convolutional layers to extract feature interactions.
5. Use pooling layers to reduce dimensions and avoid overfitting.
6. Train and Evaluate. Train the CNN model on the reshaped data.
7. Evaluate using metrics like AUC to compare its performance with the fully connected model.

```
In [59]: from tensorflow.keras.layers import Conv2D, Flatten, MaxPooling2D, Input
from sklearn.preprocessing import StandardScaler


# Reshape data for CNN
def reshape_for_cnn(X):
    # Assuming 30 features can be reshaped into 5x6 grid
    return X.values.reshape(-1, 5, 6, 1) # Add channel dimension for grayscale


# Reshape training and testing sets
X_train_cnn = reshape_for_cnn(X_train)
X_test_cnn = reshape_for_cnn(X_test)


# Define a simple CNN model
def build_cnn_model(input_shape):
    model = Sequential([
        Input(shape=input_shape),
        Conv2D(32, (3, 3), activation='relu'),
        MaxPooling2D((2, 2)),
        Dropout(0.5),
        Flatten(),
        Dense(64, activation='relu'),
        Dense(1, activation='sigmoid')
    ])
    model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['AUC'])
    return model


# Build and compile the CNN model
cnn_model = build_cnn_model((5, 6, 1))


# Train the CNN model
history_cnn = cnn_model.fit(
    X_train_cnn, y_train,
    validation_split=0.2,
    epochs=20,
    batch_size=512,
    callbacks=[tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=8,
    verbose=1
    )
```


Epoch 1/20
357/357  3s 4ms/step - AUC: 0.2867 - loss: 0.1324 - val_AUC: 0.8821 - val_loss: 0.0075


Epoch 2/20
357/357  2s 4ms/step - AUC: 0.8214 - loss: 0.0098 - val_AUC: 0.8969 - val_loss: 0.0053


Epoch 3/20
357/357  1s 4ms/step - AUC: 0.8788 - loss: 0.0064 - val_AUC: 0.9084 - val_loss: 0.0047


Epoch 4/20
357/357  1s 4ms/step - AUC: 0.8971 - loss: 0.0053 - val_AUC: 0.9092 - val_loss: 0.0044


Epoch 5/20
357/357  1s 4ms/step - AUC: 0.9271 - loss: 0.0045 - val_AUC: 0.9209 - val_loss: 0.0039


Epoch 6/20
357/357  2s 4ms/step - AUC: 0.9135 - loss: 0.0047 - val_AUC: 0.9093 - val_loss: 0.0042


Epoch 7/20
357/357  2s 4ms/step - AUC: 0.9038 - loss: 0.0050 - val_AUC: 0.9213 - val_loss: 0.0040


Epoch 8/20
357/357  2s 5ms/step - AUC: 0.9061 - loss: 0.0047 - val_AUC: 0.9215 - val_loss: 0.0041


Epoch 9/20
357/357  2s 4ms/step - AUC: 0.9394 - loss: 0.0041 - val_AUC: 0.9214 - val_loss: 0.0038


Epoch 10/20
357/357  1s 4ms/step - AUC: 0.9256 - loss: 0.0038 - val_AUC: 0.9214 - val_loss: 0.0037


Epoch 11/20
357/357  1s 4ms/step - AUC: 0.9226 - loss: 0.0044 - val_AUC: 0.9215 - val_loss: 0.0038


Epoch 12/20
357/357  2s 4ms/step - AUC: 0.9177 - loss: 0.0038 - val_AUC: 0.9211 - val_loss: 0.0036


Epoch 13/20
357/357  1s 4ms/step - AUC: 0.9124 - loss: 0.0042 - val_AUC: 0.9213 - val_loss: 0.0035


Epoch 14/20
357/357  1s 3ms/step - AUC: 0.9263 - loss: 0.0033 - val_AUC: 0.9215 - val_loss: 0.0036


Epoch 15/20
357/357  1s 3ms/step - AUC: 0.9181 - loss: 0.0040 - val_AUC: 0.9214 - val_loss: 0.0036

Epoch 16/20
357/357  2s 4ms/step - AUC: 0.9193 - loss: 0.0033 - val_AUC: 0.9215 - val_loss: 0.0036

Epoch 17/20
357/357  2s 4ms/step - AUC: 0.9490 - loss: 0.0028 - val_AUC: 0.9272 - val_loss: 0.0035

Epoch 18/20
357/357  2s 4ms/step - AUC: 0.9417 - loss: 0.0033 - val_AUC: 0.9271 - val_loss: 0.0034

Epoch 19/20
357/357  1s 4ms/step - AUC: 0.9495 - loss: 0.0030 - val_AUC: 0.9214 - val_loss: 0.0034

Epoch 20/20
357/357  2s 5ms/step - AUC: 0.9328 - loss: 0.0033 - val_AUC: 0.9275 - val_loss: 0.0034

Key Findings and Insights

Model evaluation

In [53]: *# Evaluate the simple model*

```
model_eval = simple_model.evaluate(X_test, y_test, verbose=1)
print(f"Test AUC (Simple Model): {model_eval[1]}")
```

1781/1781 ————— 2s 850us/step - AUC: 0.9509 - loss: 0.0026
Test AUC (Simple Model): 0.9484700560569763

In [56]: *# Evaluate the complex model*

```
complex_eval = complex_model.evaluate(X_test, y_test, verbose=1)
print(f"Test AUC (Complex Model): {complex_eval[1]}")
```

1781/1781 ————— 2s 878us/step - AUC: 0.9511 - loss: 0.0022
Test AUC (Complex Model): 0.9536031484603882

In [60]: *# Evaluate the CNN model*

```
cnn_eval = cnn_model.evaluate(X_test_cnn, y_test, verbose=1)
print(f"Test AUC: {cnn_eval[1]}")
```

1781/1781 ————— 2s 1ms/step - AUC: 0.9507 - loss: 0.0021
Test AUC: 0.9436067342758179

In [63]: **import** matplotlib.pyplot **as** plt

Plot Loss trends for the three models

```
plt.figure(figsize=(10, 6))
```

Simple Model

```
plt.plot(history.history['loss'], label='Simple Model - Train Loss', linestyle='solid')
plt.plot(history.history['val_loss'], label='Simple Model - Validation Loss')
```

Complex Model


```
plt.plot(history_complex.history['loss'], label='Complex Model - Train Loss', linestyle='solid')
plt.plot(history_complex.history['val_loss'], label='Complex Model - Validation Loss')
```

CNN Model

```
plt.plot(history_cnn.history['loss'], label='CNN Model - Train Loss', linestyle='solid')
plt.plot(history_cnn.history['val_loss'], label='CNN Model - Validation Loss')
```

Add plot details

```
plt.title('Loss Function Trends for All Models')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.grid()
plt.show()
```



Final Comments

The Complex Model (AUC: 0.9536) performed slightly better than the Simple Model (AUC: 0.94847), demonstrating that adding more layers and nodes helped capture more intricate patterns in the data.

The CNN Model (AUC: 0.9436), while slightly lower in performance, still shows competitive results. It is likely effective at capturing feature interactions due to its convolutional layers, but the spatial structure used in reshaping features may not have been optimal.

However, despite the high AUC, the class imbalance remains a significant challenge. The model's performance on metrics like recall for the minority (fraud) class should be closely monitored to ensure fraud detection is effective.

The performance improvement is incremental. Testing advanced techniques like hyperparameter optimization, deeper CNNs, or ensemble methods could provide more gains.

Next Steps

Further improvements could include:

1. **Using SMOTE (Synthetic Minority Oversampling Technique)** or similar methods to address the class imbalance.
2. **Hyperparameter tuning** to further optimize the model's performance.
3. Experimenting with **advanced architectures**, such as **Recurrent Neural Networks (RNNs)** to incorporate time-dependency in the data.

By addressing these areas, the model's ability to detect fraudulent transactions could be further enhanced, potentially improving both recall and precision for the minority class.